

Using Claude with Purpose

Judgment at the Keyboard

A Working Guide for Executives and Senior Leaders

Beyond setup and features — practical patterns, honest weaknesses, and a framework for getting real value from AI at the executive level

DeWayne L. Searcy

Using Claude with Purpose

Judgment at the Keyboard — A Working Guide for Executives and Senior Leaders

Copyright © 2026 DeWayne L. Searcy

All rights reserved.

Published by 12 to KNOW, LLC

First Edition, 2026

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

Claude is a product of Anthropic. "Claude" is referenced descriptively throughout this guide and is not affiliated with or endorsed by Anthropic. All product names, logos, and brands referenced herein are the property of their respective owners.

The content of this guide reflects the author's observations and experience. It is not legal, financial, medical, or professional advice. Readers should consult qualified professionals before acting on any guidance herein.

Why This Guide Exists

Most material written about using Claude — or any large language model — falls into one of two buckets.

The first is academic. It examines AI adoption at the organizational level, debates the paradox between automation and augmentation, and produces frameworks for workforce strategy. It is valuable for scholars and chief strategy officers. It does not help the executive sitting at the keyboard on Monday morning trying to get useful output.

The second is technical. It explains how to install the tool, how to configure Projects, how to connect MCP servers, how to write a good prompt if you are a developer. It is valuable for builders. It assumes a technical fluency most senior leaders do not have and does not need to acquire.

This guide occupies the space in between. It is written for the executive who is not a technologist and not a scholar — someone who is running a business, leading a team, making decisions with real consequences, and using AI as part of that work. It is not about features. It is about judgment: how to ask, how to iterate, when to trust the output, when to walk away, and where the line is between collaboration and over-reliance.

The framework below answers three questions a busy leader should be able to answer quickly:

- What is Claude actually good at, and what is it not?
- How do I ask in a way that gets real value instead of a generic answer?
- Where is the line between collaboration and over-reliance?

Every pattern, strength, and weakness described here is drawn from real executive use. No hypotheticals, no feature tours, no vendor positioning.

The question is not whether AI can help you. It is whether you will use it in ways that compound value or in ways that slowly erode the judgment you were hired for.

Part 1 — How Executives Actually Use Claude

Serious executive use falls into five categories. Each one has a different success pattern and naming the category up front changes how you should prompt and what you should expect back.

1.1 The Five Categories

Category	What It Looks Like	Typical Session
Build deliverables	Policies, training modules, spreadsheets with real formulas, branded PDFs, interactive web pages, presentation decks	Multi-turn, iterative
Strategic analysis	Market evaluations, site comparisons, competitive landscapes, scenario planning, pre-decision memos	Extended, research-heavy
Operational problem solving	Root-cause diagnosis on recurring equipment or process failures, compliance corrections, workflow troubleshooting	Focused, diagnostic
Research and context	Understanding a regulation, a market term, a person's background, an industry shift, a historical parallel	Short, single-question
Personal and family	Health questions framed for deeper conversation with a professional, genealogy, practical life decisions	Varies

A healthy mix includes all five. The five-minute practical questions do not compete with the five-hour strategy work; they are different uses of the same tool, and both are legitimate.

1.2 What Strong Usage Looks Like

Treating Claude as a collaborator, not a vending machine

The strongest users push back when something feels off. They correct the assistant when it makes an assumption that does not match their business. They flag layout problems in draft deliverables rather than accepting them. They say "that's not quite right — here is what you're missing." This is the single most valuable habit.

Building toward deployable artifacts, not summaries

Strong users finish things. A policy that gets distributed. A training module that gets launched. A spreadsheet the team actually uses. They are not generating content for its own sake; they are shipping. Summaries without a destination rarely create value.

Mixing domains without losing context

Many users compartmentalize AI use — business over here, personal over there. Strong users do not. In a single week they may move between a strategic analysis, a brand identity question, a family decision, and a technical troubleshooting session. The assistant's memory holds enough thread across sessions to keep this coherent. That range is a real advantage.

Asking second-order questions

Most users ask the assistant to do the work. The most effective users also ask the assistant to pressure-test the work. "What am I missing in this analysis?" "Where is this reasoning weakest?" "Am I over-relying on this tool?" These questions keep output honest and keep the user sharp.

Part 2 — Where Executive Usage Is Strongest

Strength 1: Context density

Output quality is proportional to the specificity of what you share. "Help me think about expansion" produces a generic summary. "Here are the zoning rules, the email from the city planner, the competitive landscape, and the internal constraint that we have never exited a market we entered — give me a go/no-go with what I am missing" produces a decision memo. The difference is not the assistant. The difference is what you put in.

The lesson: treat the first prompt like a briefing, not a question. Load it with the situation, the constraints, the history, and the goal. Everything downstream gets better.

Strength 2: Iteration stamina

First drafts are roughly 70 to 80 percent of the way to usable. The last 20 to 30 percent — the color matching, the layout correction, the tone adjustment, the language that sounds like your organization — is what separates an internal artifact from something you would put in front of the Board. Strong users do that last-mile work. Most users quit at 80 percent and wonder why the output feels generic.

The lesson: budget for three to five rounds of refinement on anything that will leave your desk. If you are not willing to iterate, you are not ready to ship.

Strength 3: Using the tool to stress-test your own thinking

The highest-leverage use of an AI assistant at the executive level is not generating content. It is pressure-testing decisions before they become expensive. "What decision factors am I missing?" "Is this the right frame?" "What would a contrarian say?" These thinking-partner questions produce the kind of value that a staff of analysts used to provide, at a fraction of the cost and time.

The lesson: the assistant's best use is often not to write the memo but to tear it apart before anyone else sees it.

Strength 4: Bridging judgment and speed

The discipline that makes AI useful at a senior level rather than dangerous is knowing what the assistant does not know. Industry-specific nuance. Internal political realities. What language actually works with a particular customer or Board member. The assistant produces fluent, plausible output in seconds. Only a human with real context can tell whether that output will survive contact with the people who matter.

Augmentation is when the assistant produces something you were going to write anyway, only faster. Outsourcing is when it produces something you could not have produced yourself and could not defend if questioned. Stay on the augmentation side.

Part 3 — Where There Is Room to Improve

These are not failures of the tool. They are common patterns where users leave value on the table or create risk for themselves.

Opportunity 1: Landing the artifacts you generate

The most common executive failure mode with AI is not overuse. It is under-execution. A library of near-finished policies, half-deployed tools, comprehensive analyses without a decision, and draft communications that never get sent. The assistant has no idea what has shipped versus what sits in your downloads folder. You do.

Practical improvement: before starting a new build, spend two minutes listing the artifacts currently in draft and which one is most valuable to finish. Use the assistant to plan the final 20 percent, not to kick off more 80-percent drafts.

Opportunity 2: Using the wrong tool for the question

Some short factual questions — a definition, a date, a basic spec — get answered fine, but the conversational overhead is not worth it over a focused web search. This matters only when many such questions pull you out of deeper work and fragment your attention.

Practical improvement: a rough mental filter. If the question has one obviously correct answer and you would accept the first search result, use a search engine. If the question has tradeoffs, requires synthesis, or benefits from your context being known, use the assistant.

Opportunity 3: The subject-matter-expert validation gap

The assistant will produce language that sounds authoritative on subjects where real domain experts exist inside your organization. Sales objection handling. Clinical language. Legal standards. Technical operating procedures. This output can be directionally right and specifically wrong — wrong enough to matter when it reaches a customer, a patient, a regulator, or a court.

Practical improvement: for any assistant-generated artifact that will be used by employees in a customer-facing, clinical, or legal context, mark it visibly as "Draft — pending SME validation." Remove the banner only after a named internal expert has reviewed. This protects the tool's credibility and yours.

Opportunity 4: Time-sensitive facts

Product naming changes, recent pricing shifts, late-breaking regulatory decisions, and fast-moving news stories — the assistant will often search the web for these, but the results can be thin, conflicting, or outdated within days. An answer that was right last week may be wrong today.

Practical improvement: for time-sensitive business decisions, treat the assistant's web output as a starting inventory, not a final answer. Verify with a primary source — the regulator's website, the vendor's pricing page, the company's own filing — before acting.

Opportunity 5: Sensitive judgment calls

Health decisions. End-of-life planning. Employment terminations. High-stakes family matters. The assistant can provide useful framing and vocabulary, but cannot replace the professional who sees the patient, the attorney who knows the jurisdiction, or the person sitting across from you who deserves your full attention.

Practical improvement: use the assistant for context-gathering and question-drafting, not for directives. Read the framing, then talk to the actual professional. Never outsource the part that requires your presence.

Part 4 — The Executive How-To Framework

This is the practical guidance layer — a framework any senior leader can pick up and use tomorrow.

4.1 The Four Question Types

Every conversation with the assistant falls into one of four question types. Naming the type before you start changes how you prompt and what you should expect back.

Question Type	What It Looks Like	How to Prompt
Build	Create a training module, a policy, a spreadsheet, a branded document	Provide business context, brand constraints, audience, and a reference for format. Expect 2 to 5 iteration cycles.
Analyze	Help me think through this decision, market, site, or risk	Share the specifics — documents, emails, constraints, history. Ask for the analysis AND for what is missing. Expect the assistant to flag gaps.
Diagnose	Why is X happening, breaking, or underperforming	Give symptoms, models, error messages, context. Ask for a ranked list of likely causes with cost-to-check each.
Research	What is, when did, who is, how does	Short and direct. Verify with a primary source if consequential.

4.2 The Context Ladder

Output quality scales with context density. A ladder, from weakest to strongest prompt:

- **Rung 1 — No context:** "Help me with expansion."
- **Rung 2 — Domain:** "Help me think about expansion for a regional services company."
- **Rung 3 — Situation:** "We are evaluating a first entry into a neighboring state. Here is the parcel, the zoning code, the competitive map."
- **Rung 4 — Specifics:** "Here is everything on Rung 3 plus our internal constraints, the email from the city planner, and the three competitors we care about. What am I missing in the go/no-go?"

Most first-time users prompt at Rung 1 or 2. Effective executive prompts are at Rung 4. The jump from Rung 2 to Rung 4 is where executive-level output starts to appear.

4.3 The Three Checks Before Acting on Output

- **Source check.** For factual claims, is this from the assistant's training, from a web search it ran, or from a document you provided? All three can be wrong in different ways. For current facts, web search beats training. For your own business, the documents you provide beat either.
- **SME check.** For organization-facing artifacts, who inside the company should validate this before it circulates? Name the person. Do not circulate first and hope.
- **Reversibility check.** If the assistant is wrong here, what does it cost? Scale your review intensity to the cost of being wrong. A personal journaling tool — low cost, low review. A compliance procedure — high cost, high review.

4.4 When Not to Use the Assistant

- **Final legal or tax opinions.** Use the assistant to understand the framework, then go to counsel.
- **Clinical diagnosis for yourself or a family member.** Use it for vocabulary and questions to ask. Then see the provider.
- **Confidential data you cannot control.** Follow your organization's AI use policy. If you do not have one, stop and write one.
- **Decisions that require your authority to own.** A Board recommendation can be informed by the assistant but must be yours.
- **Work that requires accountability to a real human.** A condolence note. A performance conversation. A termination discussion. The assistant can draft; you must own the words.

4.5 The Presence Test

The Three Checks in 4.3 test the output. This one tests you.

The augmentation-versus-outsourcing distinction (Part 2, Strength 4) gives you an output test: could I have written this myself? It is a good test, and it is retrospective — it catches work that is beyond your capability. It does not catch a subtler failure: producing work you were fully capable of producing, while not actually engaged in producing it. The artifact comes out clean. You could have written it. And yet you were absent while it was made, so if someone challenged a specific figure, you could not defend it without going back to re-read. The output passes every quality check and you have still outsourced the thinking.

So the output test needs a companion. Where the output test asks a retrospective question, the presence test asks a present-tense one:

Am I tracking this well enough, right now, that if someone challenged a specific number in the next five minutes, I could answer it without re-reading?

Both must pass. The output test is about the artifact; the presence test is about you. Note the tie to 4.4: presence there is a gate — the thing that tells you to close the laptop. Here it is a dial that must stay on even inside the work you are allowed to use the tool for.

Two things make this check easy to skip:

- **It is invisible.** Disengagement produces no error message. The work looks finished, so nothing signals that you checked out while it was made. You have to check for presence deliberately; it never announces its own absence.
- **It is worst where you are strongest.** In an unfamiliar domain you stay alert, because you know you cannot verify the output. In your home domain you relax, because coasting feels safe. It is the worst place to coast: your hard-won pattern recognition is the highest-value input in the exercise and the one thing the assistant cannot supply.

The lesson: your risk of disengagement is highest, not lowest, in the domain you know best. Expertise feels like permission to coast. It is the opposite.

The fix is a forcing function, not more willpower — willpower loses to this by design. Before you read the assistant's version of a result, write down what you expect it to show: your read, your priors, the number you would bet on. Then compare. The divergences are the point — either the tool caught something you missed, or you caught something it got wrong, and either way you are back in the loop, because you now have a stake in a prediction. Passive monitoring lulls you; active prediction keeps you sharp.

The lesson: predict before you review. If you will not commit to a prior, you are not engaged enough to trust your own approval.

Part 5 — Adopting This in Your Organization

5.1 A Layered Rollout

If this material is going to become a training or reference asset, it layers well for different audiences.

Layer 1 — Executive brief

Parts 3 and 4 only. Three to four pages. For senior leaders who will use the assistant but do not need the full picture. Pair with your AI use policy.

Layer 2 — Training module

Parts 1, 3, and 4. Delivered as a short interactive web module or a 30-minute session. Include three to four worked examples drawn from your own organization, with any sensitive details scrubbed.

Layer 3 — Educational content for a wider audience

Parts 2, 3, and 4. Written for professionals who are learning to use AI credibly. The angle is how a senior leader actually uses the tool in practice — which is content very few voices in the space are producing with real evidence behind it.

5.2 A Final Word on Over-Reliance

I will end with something that happened to me, because it is the most honest illustration I can offer of everything above.

I opened two spreadsheets — a rental branch I had analyzed for a decade, and a fresh year of the same data — with the assistant running alongside, and asked it to repeat an analysis I know cold. It did. Somewhere in the middle I realized I was staring into space. The work was getting done. I was not in it. When I looked at the finished output, I could not honestly say I understood it well enough to make the personnel decisions it pointed toward — decisions about real people.

Nothing about the output was wrong. That is what made it dangerous.

By the test in this guide, I was fine. Augmentation is when the assistant produces something you were going to write anyway, only faster; outsourcing is when it produces something you could not have produced yourself and could not defend if questioned. I could have written that analysis in my sleep. And I outsourced it anyway — because I was absent while it was made. That is the gap the presence test (4.5) exists to close: the output can pass every quality check while you are checked out, and that is still outsourcing.

The detail that stays with me is that I disengaged because I was the expert. Ten years of judgment felt like permission to coast — on the one exercise where my judgment was the entire point. That is the real shape of the risk at the senior level. It is not overuse in volume. It is the quiet handoff of thinking you are supposed to be doing, and the person most exposed to it is not the novice fumbling with the tool. It is the expert who knows the work so well they stop watching it.

So the test is not only whether you could have written it. It is whether you were there when it was written. Stay present, predict before you review, and keep the judgment you were hired for. Drift into absence — however fluent the output — and no amount of productivity gain will compensate for the judgment you let atrophy.